

Best Team Ever





Zillow Kaggle Competition

Patrick Masi-Phelps | Katie Critelli
Ningxi Xu | John Merrick

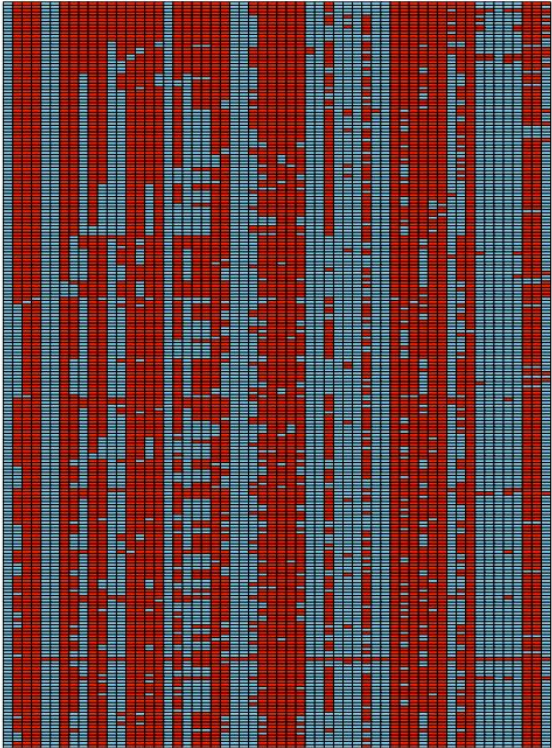
EDA: Visualization of Missingness

Missingness Plot




ngclassypeid
sqarefeet13
storytypeid
basementsqft
buildingsqft26
fireplaceflag
ruralstyletypeid
rfructiontypeid
rsquarefeet16
decktypeid
poolsize
pooltypeid10
pooltypeid2
linquencyyear
ishotlubor spa
buildingsqft17
sqarefeet15
sqarefeet50
or1squarefeet
fireplacecent
uaterbathnbr
pooltypeid7
poolcent
mberofstories
litioningtypeid
aragetotalsqft
garagecarcnt
neighborhood
rsystemtypeid
qualitytypeid
rtyzoningdesc
unitcnt
izesquarefeet
fsquarefeet12
regionidcity
fullbathcnt
ulatedbathnbr
yearbuilt
bedsquarefeet
fractandblock
valuedollarcent
regionidzip
taxamount
valuedollarcent
valuedollarcent
yandusecode
ansactiondate
logerror
sessmentyear
roomcnt
regionidcounty
fractandblock
landusetypeid
longitude
latitude
fips
bedroomcnt
bathroomcnt
parcelid

Combinations




parcelid
decktypeid
fips
latitude
regionidzip
unitcnt
taxamount

EDA: Correlation of Numeric Variables



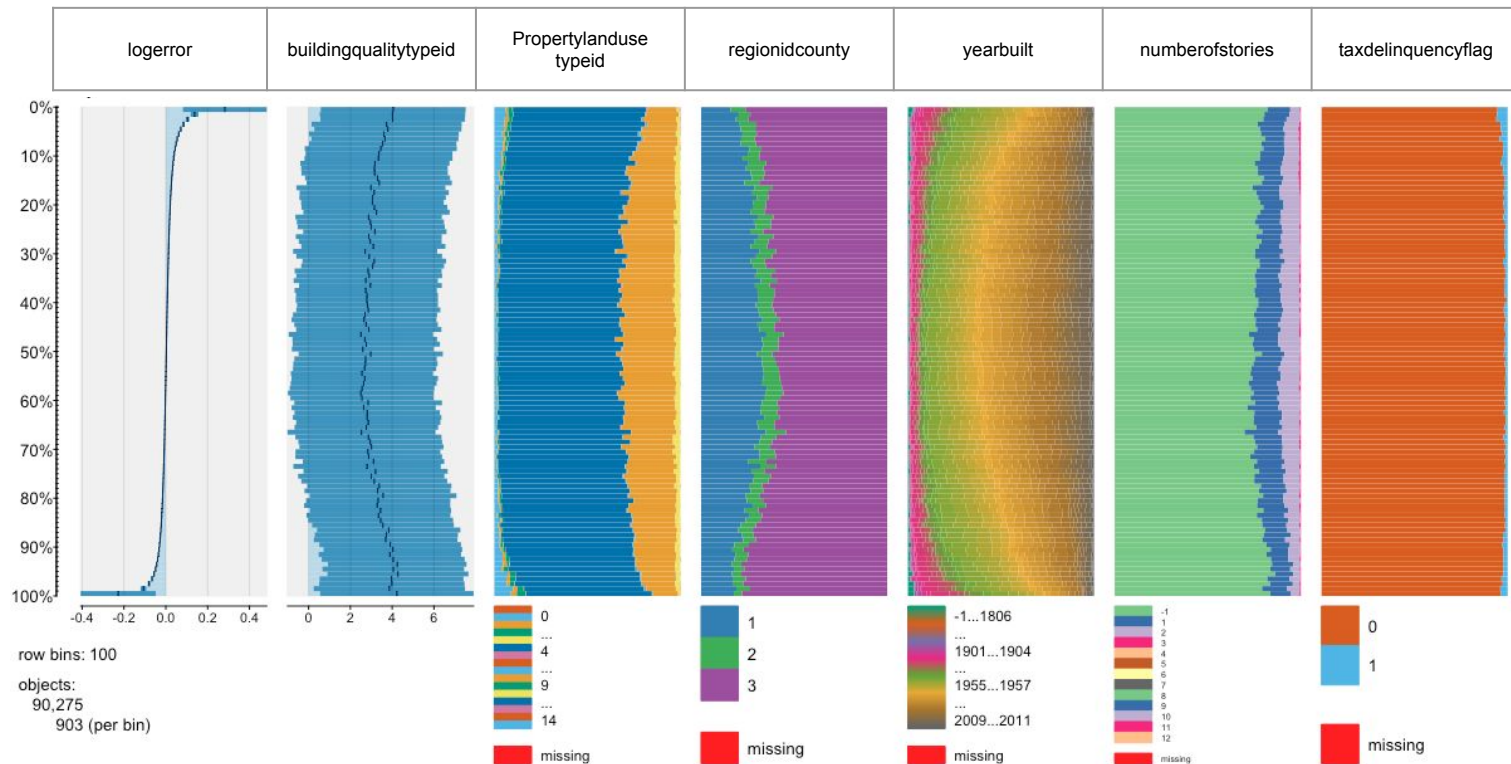
EDA: Tableplot Visualizations

- Predictor values sorted by logerror



EDA: Tableplot Visualizations

- Predictor values sorted by logerror



Data Cleaning and Imputation

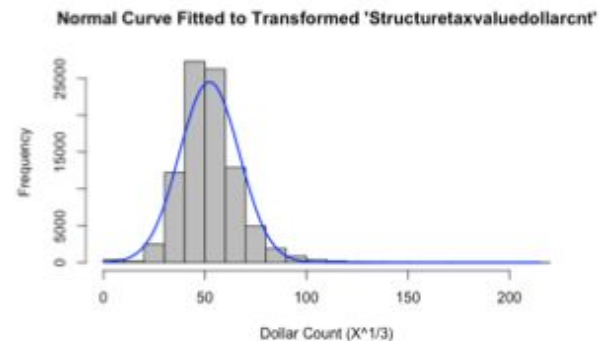



Data Cleaning and Imputation

- Does the property contain the item classified by the variable?
 - Assumed NA means this property does not contain the item: imputed NAs with 0
 - Taxdelinquencyyear, taxdelinquencyflag, lotsizesquarefeet, fireplacecnt, garagecarcnt, garagetotalsqft, poolcnt, pooltypeid7, pooltypeid2, poolsizesum, basementsqft, threequarterbathnbr, hashottuborspa
 - Imputed heatingorsystemtypeid with 13 and airconditioningtypeid with 5 (equivalent to none)
- Reason for NA is unknown: imputed with -1 per Shu's lecture code
 - Preserving any reasons for missingness that could be inherent in these particular observations
 - Unitcnt, bedroomcnt, bathroomcnt, fullbathcnt, buildingqualitytypeid, numberofstories, yearbuilt
- Other methods
 - Calculatedfinishedsquarefeet: mean imputation
 - Dropped other floor space area variables as redundant/overly missing
 - Region variables: random imputation
 - Small number of missing data points

Data Cleaning and Imputation



- Missing tax building and land assessment values
 - Examine property taxes paid
 - **23,887/23,902** observations had zero bathrooms and zero bedrooms (not missing)
 - **Likely no building**
 - Divide property taxes paid by median property tax rate (across all properties)
 - Impute this value for land assessment value
 - Impute zero for building assessment value
 - For properties with no taxation values
 - Impute average building and land assessment values
 - Group by zip code, num bed, num bath

Data Transformations



Data Transformations



- Flexible technique
 - Measured skewness of each variable
 - Found a constant C for which the skewness of the transformed function, $\log(X + C)$, was minimized



Models

GLM process trial and error

- **Fitting lr1:** $\text{logerror} \sim \text{bathroomcnt} + \text{bedroomcnt} + \text{buildingqualitytypeid} + \text{calculatedfinishedsquarefeet} + \text{fireplacecnt} + \text{fullbathcnt} + \text{garagecarcnt} + \text{lotsizesquarefeet} + \text{poolcnt} + \text{taxdelinquencyflag} + \text{structuretaxvaluedollarcnt} + \text{unitcnt} + \text{yearbuilt}, \text{numberofstories} + \text{landtaxvaluedollarcnt}$
 - **Result:** poor predictive power and several features with high VIFs (including garagecarcnt, poolcnt, etc)
- **Remedy #1:** Fit lr2, leaving out features with both high VIFs and high p-values
 - **Result:** R-squared still low, violated linear model assumptions (see selected graphs) →
- **Remedy #2:** Fit lr3, discarding additional features likely to be multicollinear with others



OLS to GLM with coefficient shrinkage

- R-squared still extremely low in lr3 after dropping majority of features (0.002606):
 - OLS likely not optimal model, because linear model assumptions all appeared to be in violation
 - P-value of F-test and all VIFs are significant
 - We do not believe that there is a legitimate linear relationship between logerror and variables
- As a result, we **fit a GLM model with regularization**
 - Fit 100 ridge and lasso models to improve model accuracy
 - Used lambda that yielded the lowest MSE
 - Reduced features used to only 6
 - Significant shrinkage in coefficients

Multiple Linear Regression with Regularization

Backward AIC Results Summary

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.433e-15	3.326e-03	0.000	1.000000	
calculatedfinishedsquarefeet.scaled	2.219e-02	5.519e-03	4.020	5.82e-05	***
lotssquarefeet.scaled	1.209e-02	3.535e-03	3.420	0.000627	***
structuretaxvaluedollarcnt.scaled	1.408e-02	5.095e-03	2.764	0.005717	**
landtaxvaluedollarcnt.scaled	-2.759e-02	4.197e-03	-6.574	4.94e-11	***
bedroomcnt.scaled	1.139e-02	4.537e-03	2.510	0.012084	*
unitcnt.scaled	-1.060e-02	3.521e-03	-3.011	0.002607	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

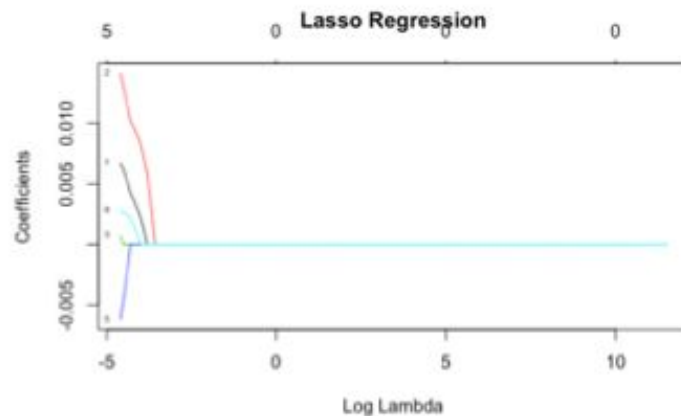
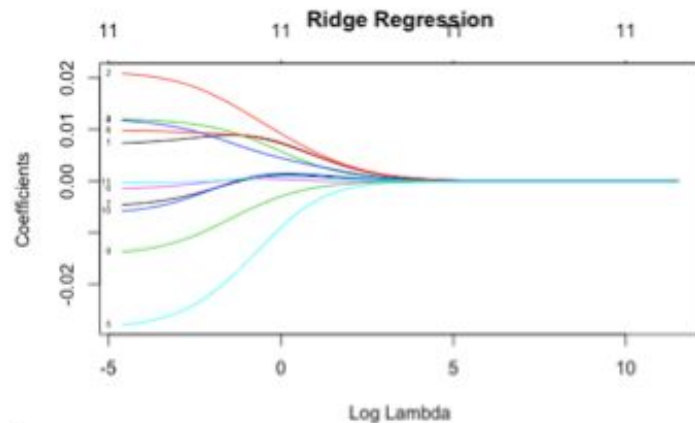
Residual standard error: 0.9992 on 90268 degrees of freedom

Multiple R-squared: 0.001577, Adjusted R-squared: 0.001511

F-statistic: 23.76 on 6 and 90268 DF, p-value: < 2.2e-16

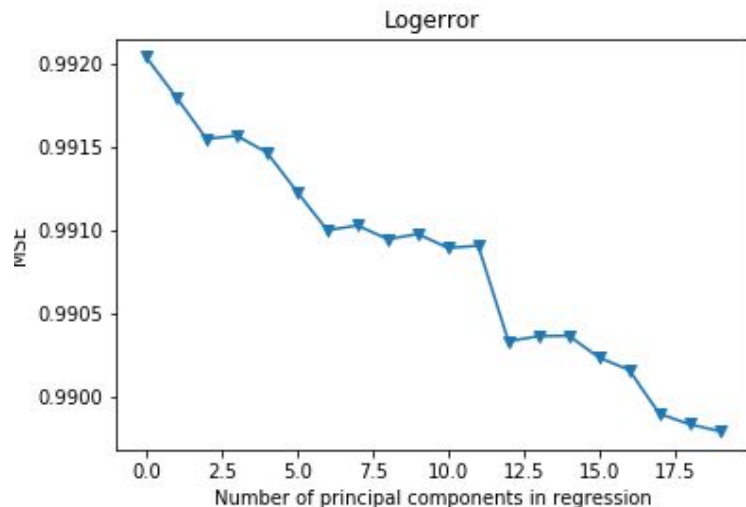
Ridge and Lasso Regression Plots

- Lasso Regression Variable Importance (lambda = .2595)
 - calculatedfinishedsquarefeet
 - bathroomcnt
 - bedroomcnt
 - landtaxvaluedollarcnt
 - lotsizesquarefeet



Principal Components Regression in Python

- Preparation for PCA
 - Scaled and centered all numeric predictors
 - LinearRegression in Scikitlearn: plotted change in MSE with each added component
 - Smallest MSE at ~17 components
 - Calculated cumulative variance explained by each added component
 - Trained regression model on training subset
 - Predicted logerror on test subset
 - Overall MSE of 1.02



Random Forest: Regression Trees


- Key tuning parameters
 - Number of randomly sampled variables considered at each split: $m_{try} = \text{variables}/3$
 - Number of trees to grow: $n_{tree} = 500$
 - Minimal size of terminal nodes: $nodesize = 5$
 - Maximum terminal nodes ($maxnodes$): if not limited by $nodesize$, will grow to max value
- Variable selection
 - Tested with all variables and gradually narrowed, eliminating less important variables
 - Broad model performed best
- Tuning and cross validation
 - Sampled 10%-25% of data set for initial broad grid searches and cross validation
 - High computational complexity of random forest calculations
 - Efficiently narrow down the search for optimal parameters
 - 75%/25% training/test split for precise model tuning

Random Forest: Regression Trees


- Optimal parameters
 - Number of variables considered at each split: 2
 - Number of trees: 1000
 - Minimal terminal node size: 12 observations
- Takeaways
 - Kaggle scoring: final RF model is worse than basic two-variable MLR!
 - MLR (logerror~yearbuilt+calculatedfinishedsquarefeet): 0.0651 MAE
 - RF: 0.0658 MAE
 - Problem in the model, or reflecting the challenges inherent in trying to predict this particular dependent variable/scoring metric?
 - Random forest model useful for assessing variable importance

Variable Importance - Random Forest

full_model



Variable Importance - Gradient Boosting



Gradient Boosting Regressor

- Numeric variables
 - Bathroomcnt, Bedroomcnt, Calculatedfinishedsquarefeet, Fireplacecnt, Fullbathcnt, Garagecarcnt, Latitude, Longitude, Lotsizesquarefeet, Poolcnt, Yearbuilt, Numberofstories, Structuretaxvaluedollarcnt, Landtaxvaluedollarcnt
- Categorical variables
 - Airconditioningtypeid, Propertylandusetypeid, Heatingorsystemtypeid, Regionidcounty, Buildingqualitytypeid, Unitcnt
- Tuning and cross validation
 - Used 5-fold grid search CV to minimize mean absolute error
- Optimal hyperparameters
 - Learning rate: .06
 - Maximum features per split: 11
 - Minimum samples split: 600
 - Subsample: .85
 - Max_depth: 11
 - Min samples per leaf: 30
 - Number of estimators: 40
- 80/20 train/test split
 - Mean Absolute Error: 0.0516168

XGBoost

- Variables - same as sklearn gradient boosting regressor (previous slide)
- Regularization, tuning and cross validation
 - Used 5-fold grid search CV to minimize mean absolute error
 - Ridge (MAE = .052879) chosen over Lasso (MAE = .052996) or Elasticnet ($\alpha = .5$) (MAE = .053169)
- Optimal hyperparameters
 - Learning rate: .06
 - Column sample by tree: .7
 - Max_depth: 7
 - Min child weight: 1
 - Number estimators: 1000
 - Subsample: .85
 - Reg alpha: 0
 - Reg lambda: 1 (Ridge)
- 80/20 train/test split
 - Mean Absolute Error: 0.0528799

Underneath the hood

Regularized Objective Function (one example)

Loss/error function

$$\mathcal{L}(\phi) = \sum_i l(\hat{y}_i, y_i) + \sum_k \Omega(f_k)$$

Complexity penalty function

$$\text{where } \Omega(f) = \gamma T + \frac{1}{2} \lambda \|w\|^2$$


“Weak learning” trees incorporate estimates from the prior tree

At the t-th iteration...

$$\mathcal{L}^{(t)} = \sum_{i=1}^n l(y_i, \hat{y}_i^{(t-1)} + f_t(\mathbf{x}_i)) + \Omega(f_t)$$


$\lambda = 1$ asserts L2 Ridge regression. L1 Lasso and ElasticNet are alternative options

Boosted Models: Logerror Distributions



K-Nearest Neighbors


- Tuning and cross validation
- Used 10-fold CV to minimize mean absolute error across K from 1-50



K-Nearest Neighbors

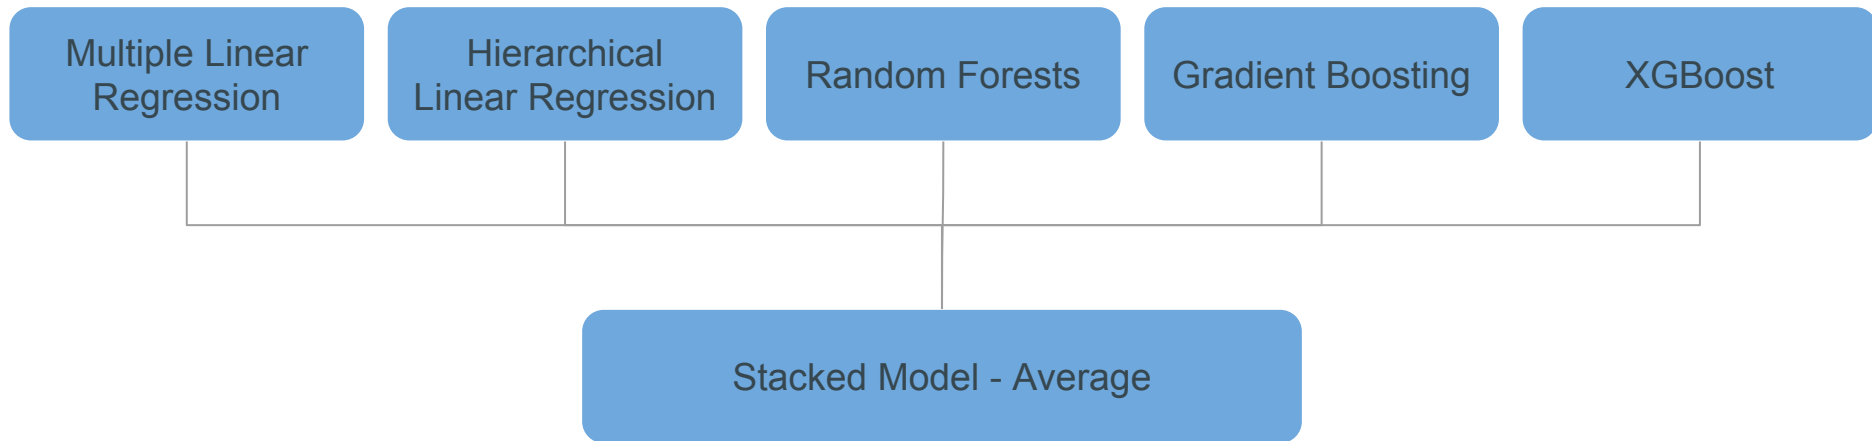
- Optimal hyperparameters:

- Weights: uniform
(rather than distance)
- K-Neighbors: 10
(near elbow)
- R^2 : .1144
- Mean Absolute Error: 0.0566785



- Run time on training/test data set (~90,000 rows): ~ 10 minutes
- Run time on full properties data set (~3,000,000 rows): **unknown (7+ hours)**
- **Has potential to be a predictive model - difficult to scale - more processing needed**

Model Stacking - Kaggle Submission



Our stacked model took the average logerror prediction across all models (equal weighting) for each property in the specified months.

Public Kaggle MAE: 0.0646177 - Rank #774

Any Questions?

Appendix: Hierarchical Linear Model

- Feature importance plots suggest that location variables are important
- HLM can highlight in-subgroup variation that is otherwise obscured
 - Treated regionidcity as the level-one subgroup
 - Coefficient of each regionidcity value was used as random variable to estimate a new linear model
 - Selected most important variables using random forest feature importance plot
 - $\text{logerror} \sim 1 + \text{structuretaxvaluedollarcnt} + \text{landtaxvaluedollarcnt} + \text{calculatedfinishedsquarefeet} + \text{lotsizesquarefeet} + \text{bedroomcnt} + \text{longitude} + \text{latitude}$
 - `groups = train['propertycountylandusecode']`
- Didn't perform well independently (MAE of 0.0660), but improved the performance of the stacked Kaggle submission