


# Best Team Ever



Zillow Kaggle Competition

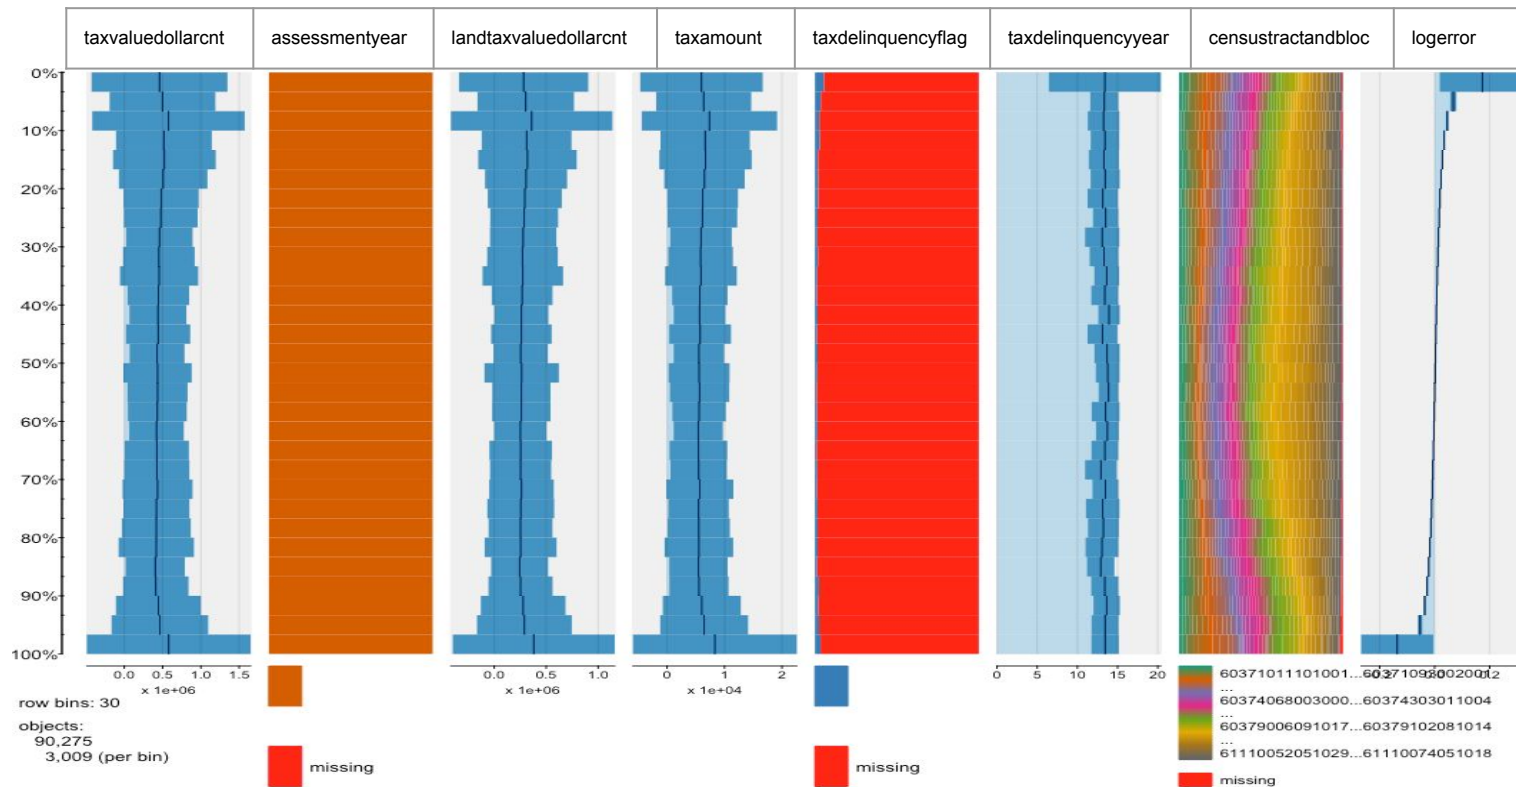
Patrick Masi-Phelps | Katie Critelli  
Ningxi Xu | John Merrick

# EDA: Correlation of Numeric Variables




# EDA: Tableplot Visualizations

- Predictor values sorted by logerror





# EDA: Tableplot Visualizations

- Predictor values sorted by logerror



[REDACTED]

## Missingness Plot



# Data Cleaning and Imputation

# Data Cleaning and Imputation

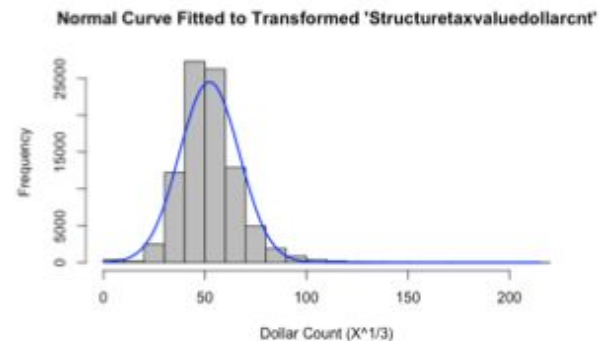



- Does the property contain the item classified by the variable?
  - Assumed NA means this property does not contain the item: imputed NAs with 0
  - Taxdelinquencyyear, taxdelinquencyflag, lotsizesquarefeet, fireplacecnt, garagecarcnt, garagetotalsqft, poolcnt, pooltypeid7, pooltypeid2, poolsizesum, basementsqft, threequarterbathnbr, hashottuborspa
  - Imputed heatingorsystemtypeid with 13 and airconditioningtypeid with 5 (equivalent to none)
- Reason for NA is unknown: imputed with -1 per Shu's lecture code
  - Preserving any reasons for missingness that could be inherent in these particular observations
  - Unitcnt, bedroomcnt, bathroomcnt, fullbathcnt, buildingqualitytypeid, numberofstories, yearbuilt
- Other methods
  - Calculatedfinishedsquarefeet: mean imputation
    - Dropped other floor space area variables as redundant/overly missing
  - Region variables: random imputation
    - Small number of missing data points

# Data Cleaning and Imputation

- Missing tax building and land assessment values
  - Examine property taxes paid
    - **23,887/23,902** observations had zero bathrooms and zero bedrooms (not missing)
      - **Likely no building**
    - Divide property taxes paid by median property tax rate (across all properties)
      - Impute this value for land assessment value
      - Impute zero for building assessment value
  - For properties with no taxation values
    - Impute average building and land assessment values
      - Group by zip code, num bed, num bath





# Data Transformations



# Data Transformations



- Flexible technique
  - Measured skewness of each variable
  - Found a constant  $C$  for which the skewness of the transformed function,  $\log(X + C)$ , was minimized



# Models

# GLM process trial and error

- **Fitting lr1:**  $\text{logerror} \sim \text{bathroomcnt} + \text{bedroomcnt} + \text{buildingqualitytypeid} + \text{calculatedfinishedsquarefeet} + \text{fireplacecnt} + \text{fullbathcnt} + \text{garagecarcnt} + \text{lotsizesquarefeet} + \text{poolcnt} + \text{taxdelinquencyflag} + \text{structuretaxvaluedollarcnt} + \text{unitcnt} + \text{yearbuilt}, \text{numberofstories} + \text{landtaxvaluedollarcnt}$ 
  - **Result:** poor predictive power and several features with high VIFs (including garagecarcnt, poolcnt, etc)
- **Remedy #1:** Fit lr2, leaving out features with both high VIFs and high p-values
  - **Result:** R-squared still low, violated linear model assumptions (see selected graphs) →
- **Remedy #2:** Fit lr3, discarding additional features likely to be multicollinear with others



# OLS to GLM with coefficient shrinkage

- R-squared still extremely low in lr3 after dropping majority of features (0.002606):
  - OLS likely not optimal model, because linear model assumptions all appeared to be in violation
  - P-value of F-test and all VIFs are significant
  - We do not believe that there is a legitimate linear relationship between logerror and variables
- As a result, we **fit a GLM model with regularization**
  - Fit 100 ridge and lasso models to improve model accuracy
  - Used lambda that yielded the lowest MSE
  - Reduced features used to only 6
  - Significant shrinkage in coefficients

# Multiple Linear Regression with Regularization

## Backward AIC Results Summary

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	3.433e-15	3.326e-03	0.000	1.000000	
calculatedfinishedsquarefeet.scaled	2.219e-02	5.519e-03	4.020	5.82e-05	***
lotssquarefeet.scaled	1.209e-02	3.535e-03	3.420	0.000627	***
structuretaxvaluedollarcnt.scaled	1.408e-02	5.095e-03	2.764	0.005717	**
landtaxvaluedollarcnt.scaled	-2.759e-02	4.197e-03	-6.574	4.94e-11	***
bedroomcnt.scaled	1.139e-02	4.537e-03	2.510	0.012084	*
unitcnt.scaled	-1.060e-02	3.521e-03	-3.011	0.002607	**

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1



Residual standard error: 0.9992 on 90268 degrees of freedom

Multiple R-squared: 0.001577, Adjusted R-squared: 0.001511

F-statistic: 23.76 on 6 and 90268 DF, p-value: < 2.2e-16

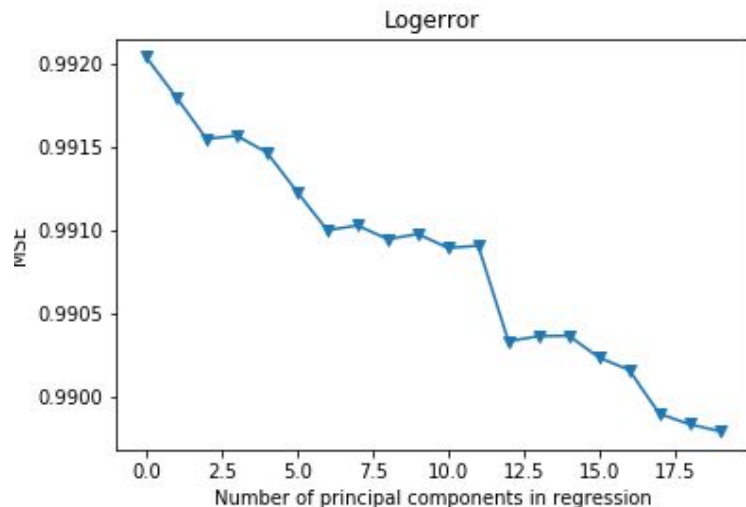
# Ridge and Lasso Regression Plots

- Lasso Regression Variable Importance (lambda = .2595)
  - calculatedfinishedsquarefeet
  - bathroomcnt
  - bedroomcnt
  - landtaxvaluedollarcnt
  - lotssquarefeet



# Principal Components Regression in Python

- Preparation for PCA
  - Scaled and centered all numeric predictors
  - LinearRegression in Scikitlearn: plotted change in MSE with each added component
  - Smallest MSE at ~17 components
  - Calculated cumulative variance explained by each added component
  - Trained regression model on training subset
  - Predicted logerror on test subset
  - Overall MSE of 1.02





# Random Forest: Regression Trees


- Key tuning parameters
  - Number of randomly sampled variables considered at each split:  $m_{try} = \text{variables}/3$
  - Number of trees to grow:  $n_{tree} = 500$
  - Minimal size of terminal nodes:  $nodesize = 5$
  - Maximum terminal nodes ( $maxnodes$ ): if not limited by  $nodesize$ , will grow to max value
- Variable selection
  - Tested with all variables and gradually narrowed, eliminating less important variables
  - Broad model performed best
- Tuning and cross validation
  - Sampled 10%-25% of data set for initial broad grid searches and cross validation
    - High computational complexity of random forest calculations
    - Efficiently narrow down the search for optimal parameters
  - 75%/25% training/test split for precise model tuning

# Random Forest: Regression Trees


- Optimal parameters
  - Number of variables considered at each split: 2
  - Number of trees: 1000
  - Minimal terminal node size: 12 observations
- Takeaways
  - Kaggle scoring: final RF model is worse than basic two-variable MLR!
    - MLR (logerror~yearbuilt+calculatedfinishedsquarefeet): 0.0651 MAE
    - RF: 0.0658 MAE
  - Problem in the model, or reflecting the challenges inherent in trying to predict this particular dependent variable/scoring metric?
  - Random forest model useful for assessing variable importance

# Variable Importance - Random Forest

full\_model



# Variable Importance - Gradient Boosting



# Gradient Boosting Regressor

- Numeric variables
  - Bathroomcnt, Bedroomcnt, Calculatedfinishedsquarefeet, Fireplacecnt, Fullbathcnt, Garagecarcnt, Latitude, Longitude, Lotsizesquarefeet, Poolcnt, Yearbuilt, Numberofstories, Structuretaxvaluedollarcnt, Landtaxvaluedollarcnt
- Categorical variables
  - Airconditioningtypeid, Propertylandusetypeid, Heatingorsystemtypeid, Regionidcounty, Buildingqualitytypeid, Unitcnt
- Tuning and cross validation
  - Used 5-fold grid search CV to minimize mean absolute error
- Optimal hyperparameters
  - Learning rate: .06
  - Maximum features per split: 11
  - Minimum samples split: 600
  - Subsample: .85
  - Max\_depth: 11
  - Min samples per leaf: 30
  - Number of estimators: 40
- 80/20 train/test split
  - Mean Absolute Error: 0.0516168

# XGBoost

- Variables - same as sklearn gradient boosting regressor (previous slide)
- Regularization, tuning and cross validation
  - Used 5-fold grid search CV to minimize mean absolute error
  - Ridge (MAE = .052879) chosen over Lasso (MAE = .052996) or Elasticnet ( $\alpha = .5$ ) (MAE = .053169)
- Optimal hyperparameters
  - Learning rate: .06
  - Column sample by tree: .7
  - Max\_depth: 7
  - Min child weight: 1
  - Number estimators: 1000
  - Subsample: .85
  - Reg alpha: 0
  - Reg lambda: 1 (Ridge)
- 80/20 train/test split
  - Mean Absolute Error: 0.0528799

## Underneath the hood

### Regularized Objective Function (one example)

Loss/error function

$$\mathcal{L}(\phi) = \sum_i l(\hat{y}_i, y_i) + \sum_k \Omega(f_k)$$

Complexity penalty function

$$\text{where } \Omega(f) = \gamma T + \frac{1}{2} \lambda \|w\|^2$$


“Weak learning” trees incorporate estimates from the prior tree

### At the t-th iteration...

$$\mathcal{L}^{(t)} = \sum_{i=1}^n l(y_i, \hat{y}_i^{(t-1)} + f_t(\mathbf{x}_i)) + \Omega(f_t)$$


$\lambda = 1$  asserts L2 Ridge regression. L1 Lasso and ElasticNet are alternative options

# Boosted Models: Logerror Distributions



# K-Nearest Neighbors

- Tuning and cross validation
- Used 10-fold CV to minimize mean absolute error across K from 1-50






# K-Nearest Neighbors

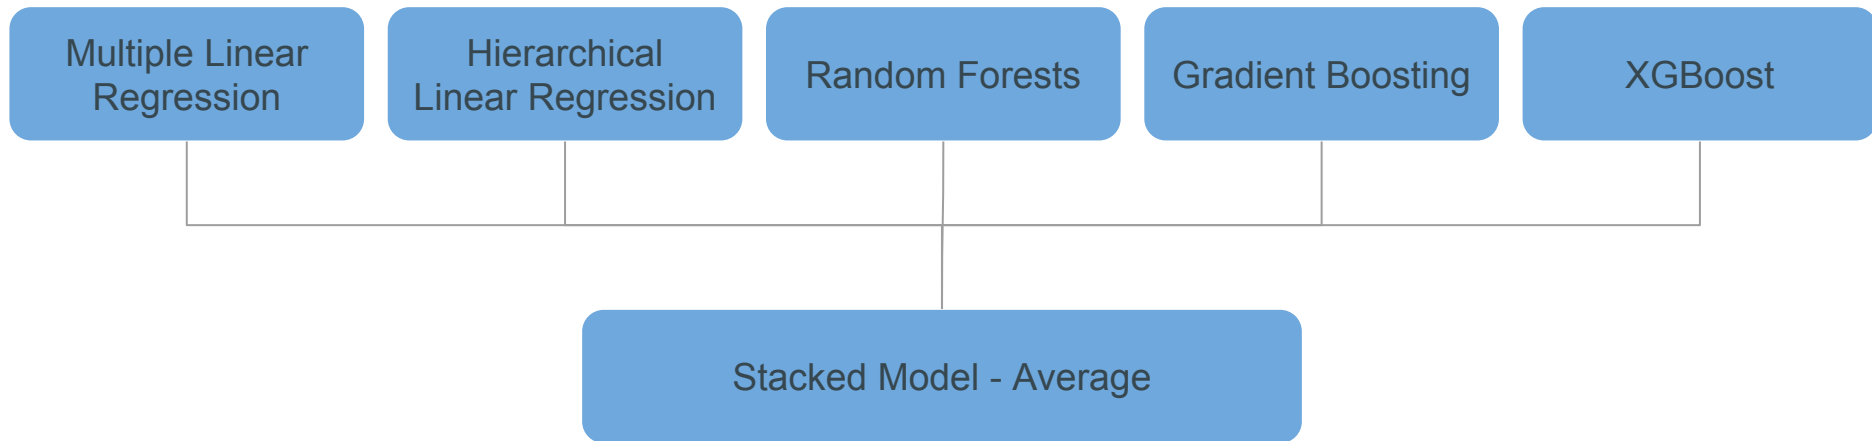
- Optimal hyperparameters:

- Weights: uniform  
(rather than distance)
- K-Neighbors: 10  
(near elbow)
- $R^2$ : .1144
- Mean Absolute Error: 0.0566785



- Run time on training/test data set (~90,000 rows): ~ 10 minutes
- Run time on full properties data set (~3,000,000 rows): **unknown (7+ hours )**
- **Has potential to be a predictive model - difficult to scale - more processing needed**

# Model Stacking - Kaggle Submission



Our stacked model took the average logerror prediction across all models (equal weighting) for each property in the specified months.

Public Kaggle MAE: 0.0646177 - Rank #774

**Any Questions?**

# Appendix: Hierarchical Linear Model

- Feature importance plots suggest that location variables are important
- HLM can highlight in-subgroup variation that is otherwise obscured
  - Treated regionidcity as the level-one subgroup
  - Coefficient of each regionidcity value was used as random variable to estimate a new linear model
  - Selected most important variables using random forest feature importance plot
    - $\text{logerror} \sim 1 + \text{structuretaxvaluedollarcnt} + \text{landtaxvaluedollarcnt} + \text{calculatedfinishedsquarefeet} + \text{lotsizesquarefeet} + \text{bedroomcnt} + \text{longitude} + \text{latitude}$
    - `groups = train['propertycountylandusecode']`
- Didn't perform well independently (MAE of 0.0660), but improved the performance of the stacked Kaggle submission